

EECE6036: Intelligent Systems  
Homework # 3

Zuguang Liu (M10291593)

November 15, 2020

# 1 Neural Network Classifier with Hidden Layer

## 1.1 Problem Statement

A classifier shall be implemented by a two-layer neural network to identify hand-written digits from MNIST dataset. The given MNIST dataset includes 5,000 labeled images, each of which writes digits ‘0’ to ‘9’ in different fashions, and is thus labeled into 10 classes. The images are represented by  $28 \times 28$  pixels of normalized greyscale color values.

## 1.2 System Description

Since neural network is a supervised learning model that requires training, there shall be an exclusive split of the dataset for training and testing. Consequently, apart from randomly shuffled, the data is also partitioned into 4,000 data points for training, and 1,000 for testing. The partitioning shall be stratified over the 10 classes, i.e., having the same number of data points for each class, in both the training set and testing set. In addition, the labels of the data are one-hot encoded to expand the output feature space to 10 dimensions.

The two-layer neural network consists of a hidden layer and an output layer. **The hidden layer incorporates 128 neurons** that transfer all 784 pixel values of an image into a 128-dimensional feature space. Then the output layer uses this feature space to classify the image with an array of 10 elements, each of which represents the probability that the image can belong to a class. Hence the input and output of the network matches the images and labels in dimensions respectively. **Both layers use the *sigmoid* activation function (1) and a weight initialization scheme known as “Xavier initialization”** showed in (2), where  $\mathcal{U}[a, b]$  is the uniform distribution in the interval  $(a, b)$ ,  $n_{in}$  is the input dimension of a layer, and  $n_{out}$  is the output dimension of a layer [1]

$$f(x) = \frac{1}{1 + e^{-x}}, \quad f'(x) = f(1 - f) \quad (1)$$

$$W \sim \mathcal{U}\left[-\sqrt{\frac{6}{n_{in} + n_{out}}}, \sqrt{\frac{6}{n_{in} + n_{out}}}\right] \quad (2)$$

The training of the model is done by back propagation repeated for multiple epochs. The training data is exclusively partitioned into 3,000 points for back propagation, and 1,000 points for validation. Within one epoch, all 3,000 points are shuffled, then used to adjust weights and biases. **The model is tested on the validation set every 10 epochs, resulting in a series of on-line training errors.** The error is calculated by  $(1 - \text{balanced accuracy})$ , where the balanced accuracy is the hit rate when comparing the true class and the predicted class using “winner-take-all” strategy over the output array. To further improve the training efficiency, several mechanisms are used in the training algorithm, including:

- a. The gradient decent has an additional term to implement momentum, demonstrated in (3), where  $\eta = 0.05$ ,  $\alpha = 0.8$ , and  $J$  is loss function implemented by  $J_2$  loss.

$$\Delta w_{ij}(t) = -\eta \frac{\partial J}{\partial w_{ij}} + \alpha \Delta w_{ij}(t - 1) \quad (3)$$

- b. **Operating thresholds of 0.25 and 0.75 are used** so that output  $\in [0, 0.25)$  is considered 0 when the corresponding truth is 0, and output  $\in (0.75, 1]$  is considered 1 when the corresponding truth is 1.
- c. Though the total repetition is 500 epochs, **an early stopping policy is used so that training stops when the on-line training loss does not improve for 50 epochs compared to the minimum loss over the whole training session.** As the validation set used for on-line testing is separate from the data used in back propagation, this ensures the model does not overfit the data.

After training, the resulting model is tested on the test dataset for analysis.

### 1.3 Results

Fig. 1 (a) and (b) are confusion matrix that demonstrates the performance of the classifier on the training and testing dataset, respectively. Fig. 2 is show the change of the on-line training error over epochs.


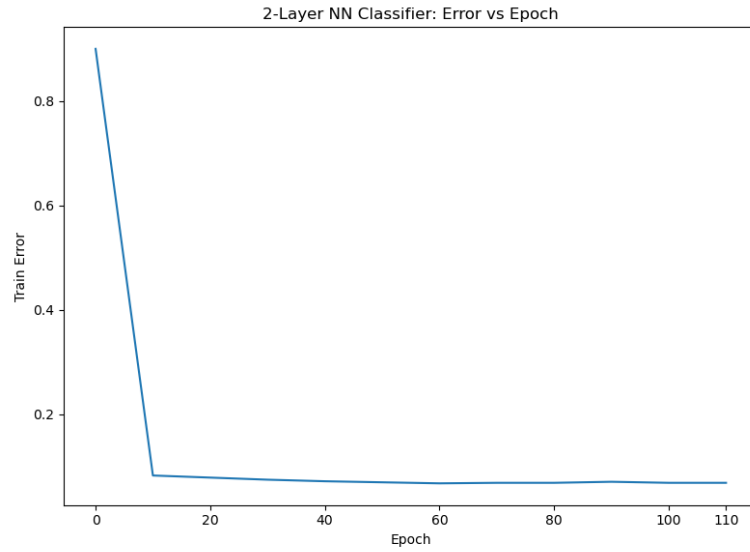


Fig. 1: Confusion matrix of the classifier on the train data (a) and test data (b).



\* Note: Training early stopped at epoch 110, then restored to epoch 60, when error is 0.068.

Fig. 2: The time series of the error (1-balanced accuracy) vs epochs.

## 1.4 Analysis of Results

The overall hit-rate on the train and test data is 98.2% and 92.5%, making the model a success. The 6.1% difference could be caused by a slight level of overfitting, but both accuracy is good in general for a classifier.

In the confusion matrices (Fig. 1), columns mark the actual class by the label, and rows mark the prediction from the model. In (b), since the test dataset are made so that there are 100 images for each ‘true class’, the numbers in each column sums up to 100, so the numbers in the diagonal is the percent hit rate for each class. Images of classes ‘0’, ‘1’, ‘3’, ‘5’, ‘6’, ‘7’ have over 90% hit rate for example, while other classes are also predicted well with the lowest hit rate at 88%. Looking back into (a), we can find there are at minimum 391 correctly classified images for every class, which is still better than the worst result on the test data. Overfitting may have affected the result, but it could also be due to the way the digits are written has too many degrees of freedom and the classifier was not able to pick up all of them.

To further improve the performance of the classifier, changes in several directions could be considered.

- a. Improvements on the quantity and quality of the data could be used. Quantitatively, the model could have larger number of data points for training and testing. Qualitatively, the images can be more defined in the resolution.
- b. The hyper-parameters of the model could be fully tested using grid search.
- c. A more advanced optimization algorithm could be used during training. For example, instead of a single term to represent the first derivative of  $\Delta w$  (3), the *ADAM* algorithm includes the second derivative also to further improve the efficiency of gradient descent.

## 2 Neural Network Autoencoder with Hidden Layer

### 2.1 Problem Statement

An auto-encoder network with one hidden layer shall be constructed for the same dataset, with the objective of reconstructing the image exactly. Therefore, the two-layer network has 784 inputs, and 784 outputs, representing a pixel in a  $28 \times 28$  image. The number of perceptrons in the hidden layer should be the same as that of the one in the classifier.

### 2.2 System Description

Since there are 128 neurons in the hidden layer of the classifier, the **autoencoder also has 128 hidden neurons**. The architecture of the autoencoder is similar to that of the classifier, except the output layer has 784 neurons to offer 784 values to regenerate the image. **Sigmoid (1) and Xavier initialization policy (2) are still used for both layers**. The training of the model also uses similar algorithm with minor modifications, including:

- The input image and the ground truth used in the training are the same per step in the back propagation.
- The back propagation uses **momentum-based gradient descent (3) with  $\eta = 0.01, \alpha = 0.8$** .
- The on-line training error is calculated by the average  $J_2$  loss function over all data points in the validation set.** This is presented by (4), where  $N$  is the number of data points,  $y_i^n$  and  $\hat{y}_i^n$  are the original and predicted  $i$ -th pixel value on the  $n$ -th image, respectively. However, the same early-stop policy is used as well, so **the training will stop after no improvement on the error for 50 epochs**.

$$\bar{J}_2 = \frac{1}{N} \sum_{n=1}^N J_2 = \frac{1}{1000} \sum_{n=1}^{1000} \left( \frac{1}{2} \sum_{i=1}^{784} (y_i^n - \hat{y}_i^n)^2 \right) \quad (4)$$

After training, the autoencoder is again tested on the test dataset for analysis.

### 2.3 Results

Fig. 3 demonstrates the performance of the autoencoder overall and at each class, where the error is calculated by (4). The on-line training error vs epochs is shown in Fig. 4.




Fig. 3: Performance of the autoencoder on the training and test set.

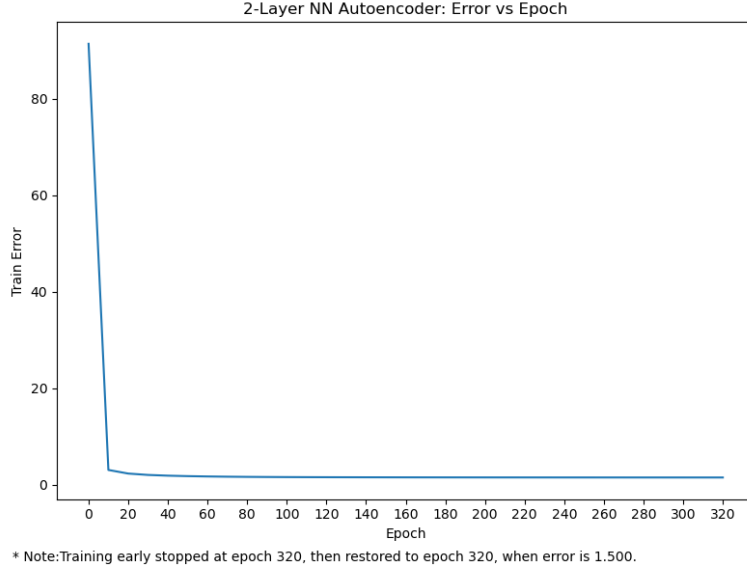


Fig. 4: On-line training error vs training epochs over time

## 2.4 Features

Weights of 20 neurons in the hidden layer of the classifier (Fig. 5) and the autoencoder (Fig. 6) are normalized and illustrated by  $28 \times 28$  images, as the feature space is 784-dimensional, in accordance to 784 pixels of the original images. Though neurons are chosen randomly, the selection of neurons in both model uses the same indexes, so that neurons in the same position are compared between the classifier and the autoencoder.

The original expectation was that the hidden features in the autoencoder could demonstrate more distinct shapes or patterns than the classifier, but the result suggests otherwise. It was surprising to see larger clusters of approximate pixel values in the classifier features than in the autoencoder features. This is discussed further in the analysis.




Fig. 5: Feature map of 20 neurons in the hidden layer of the classifier network.




Fig. 6: Feature map of 20 neurons in the hidden layer of the classifier network.

## 2.5 Sample Outputs

Fig. 7 demonstrates 8 reconstructed images (i)-(p) compared against the original images (a)-(h).

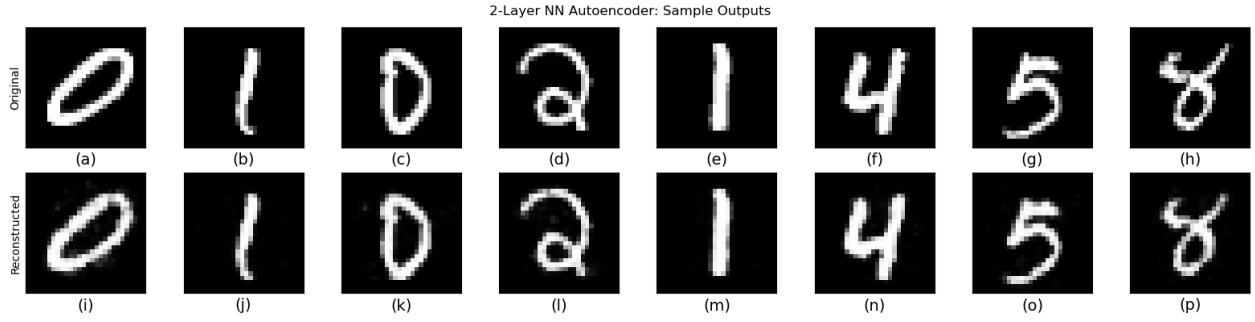


Fig. 7: 8 Sample output of the autoencoder compared to original images.

## 2.6 Analysis of Results

The resulting autoencoder is successfully trained with good level of accuracy. Mathematically, when the model is applied on the testing dataset, the overall  $J_2$  error between the output and the original pixels is around 1.6 per image (Fig. 3), similar to the end on-line training error (Fig. 4). Visually, the reconstructed images look very much alike the original images (Fig. 7). From a detailed observation, there is minor noise in (p) that writes ‘2’, (i) that writes ‘5’ and (l) that writes ‘8’. This aligns with the  $J_2$  loss in Fig. 3, as class 2, 5, and 8 have the highest errors among all classes.

Though both classifier and autoencoder network do well in their own problems, the difference in the patterns of the hidden layer feature was quite astonishing at first. It is theorized that the checkerboard pattern in the autoencoder hidden features (Fig. 6) is due to the need for the hidden layer to encode spatial correlations between pixels while reducing the feature space dimension. The model meets this goal by a set of weights that “accepts” some pixels ( $w_{ij} > \bar{w}$ ) and “denys” some pixels ( $w_{ij} < \bar{w}$ ) periodically at a certain “sampling rate”, thus the checkerboard pattern. In the classifier (Fig. 5), however, there is no need to encode local correlations, so pixels that are spatially close are allowed to be weighted similarly, forming more visual clusters. Due to limited time and knowledge constraints, this hypothesis is not verified yet but may offer a reasonable explanation.

Lastly, it is important that the training efficiently converges to such result, thanks to a well-defined model, detailed training policy (momentum-based gradient descent, early stop, etc.) and a finely-tuned set of hyper-parameters. The success of both models will provide an amount of experience in the future machine learning problems.

## References

- [1] X. Glorot and Y. Bengio, “Understanding the difficulty of training deep feedforward neural networks,” p. 8.

## Appendix A Python Code: preprocess.py

```
1  #!/usr/bin/env python3
2  # -*- coding: utf-8 -*-
3  """
4  Created on Sat Oct 24 16:13:36 2020
5
6  @author: liu
7  """
8
9  import numpy as np
10 import csv
11 import settings
12 import json
13 import pandas as pd
14
15 def prepare_img(img_a):
16     img_a = np.array(img_a).flatten()
17     img_a = 1 - img_a / np.linalg.norm(img_a)
18     return np.reshape(img_a, (-1, int(len(img_a)**0.5)), order='F')
19
20 def get_rand_list(length):
21     length = int(length)
22     return np.random.choice(length, length, replace=False)
23
24 def prepare_data():
25     x_db = []
26     with open(str(settings.X_FILE)) as csv_file:
27         csv_reader = csv.reader(csv_file, delimiter='\t')
28         for row in csv_reader:
29             x_db.append([float(x) for x in row])
30
31     y_db = []
32     with open(str(settings.Y_FILE)) as csv_file:
33         csv_reader = csv.reader(csv_file, delimiter='\t')
34         for row in csv_reader:
35             y_db.append(int(row[0]))
36
37     print('Distribution of original dataset:', np.bincount(y_db))
38
39     train_i, test_i = stratify_split(y_db,
40                                     settings.SIZES['train']/settings.SIZES['y'])
41
42     train_db = {'x': [x_db[i] for i in train_i],
43                'y': [y_db[i] for i in train_i]}
44
45     test_db = {'x': [x_db[i] for i in test_i],
46                'y': [y_db[i] for i in test_i]}
47
48     print('Distribution of train dataset:', np.bincount(train_db['y']))
49     print('Distribution of test dataset:', np.bincount(test_db['y']))
50
51     test_db['y'] = np.eye(settings.SIZES['classes'])[test_db['y']].tolist()
52     train_db['y'] = np.eye(settings.SIZES['classes'])[train_db['y']].tolist()
53
54     with open(str(settings.TRAIN_FILE), 'w') as f:
55         json.dump(train_db, f)
56
57     print("Saved train data in", settings.TRAIN_FILE)
58
59     with open(str(settings.TEST_FILE), 'w') as f:
60         json.dump(test_db, f)
61
62     print("Saved test data in", settings.TEST_FILE)
63
64 def stratify_split(y, ratio):
65     if len(np.array(y).shape) > 1: # collapse for one hot
66         y = np.argmax(y, axis=1)
67     df = pd.DataFrame(y).groupby(0) # Sort data by class
68     indxs = [] # buffer for indexes
69     for _, g in df:
70         indxs.append(g.index.to_numpy()) # indexes of each class take a row
71     indxs = np.array(indxs)
72     p1_indx = indxs[:, :int(indxs.shape[1]*ratio)].flatten() # partition 1
73     np.random.shuffle(p1_indx) # mix index
74     p2_indx = indxs[:, int(indxs.shape[1]*ratio):].flatten() # partition 2
75     np.random.shuffle(p2_indx) # mix index
76     return p1_indx, p2_indx
77
78 def get_test():
79     with open(str(settings.TEST_FILE), 'r') as f:
80         return json.load(f)
81
82 def get_train():
83     with open(str(settings.TRAIN_FILE), 'r') as f:
84         return json.load(f)
85
86 def find_ES(train_errors, max_epochs):
87     epochs = 10*np.arange(len(train_errors))
88     epochs = np.flip(epochs)
89     train_errors = np.flip(train_errors)
90     if len(train_errors) < max_epochs:
91         return "Training early stopped at epoch {}, then restored to epoch {}, when error is {:.3f}.".format(epochs[0], epochs
92         [np.argmax(train_errors)], np.min(train_errors))
93     else:
94         return "Last training error is {:.3f}.".format(train_errors[0])
95
96 if __name__ == '__main__':
97     prepare_data()
```



## Appendix B Python Code: nn.py

```
1  #!/usr/bin/env python3
2  # -*- coding: utf-8 -*-
3  """
4  Created on Sun Oct 25 14:15:46 2020
5
6  @author: liu
7  """
8  import numpy as np
9  from preprocess import get_rand_list, stratify_split
10 import json
11 from tqdm import trange
12
13
14 class DenseLayer(object): # fully connected layer
15     def __init__(self, n_input, n_neurons, activation=None, trainable=True,
16                 weights=None):
17         """
18         Initialize a fully-connected layer
19
20         Parameters
21         -----
22         n_input : uint
23             number of input nodes.
24         n_neurons : uint
25             number of neurons / output nodes.
26         activation : str, optional
27             activation function name. The default is None.
28         weights : np array, optional
29             matrix for weights. The default is None.
30
31         Returns
32         -----
33         None.
34
35         """
36
37         if weights is None:
38             a = np.sqrt(6/(n_input+n_neurons))
39             self.weights = np.random.uniform(low=-a, high=+a, size=(n_input+1, n_neurons)) #Xavier initialization
40         else:
41             weights = np.array(weights)
42             if weights.shape == (n_input+1, n_neurons):
43                 self.weights = weights
44             else:
45                 raise ValueError("Given weights does not match given dimensions")
46         self.trainable = trainable
47
48         self.last_dweights = np.zeros((n_input+1, n_neurons))
49
50         self.activation = activation
51         self.last_activation = None
52         self.error = None
53         self.delta = None
54
55     def set_trainable(self, trainable):
56         """
57         Configure if the layer is trainable
58
59         Parameters
60         -----
61         trainable : bool
62             whether the layer is trainable.
63
64         Returns
65         -----
66         None.
67
68         """
69         self.trainable = trainable
70
71     def call(self, x):
72         """
73         Calculate the output given input
74
75         Parameters
76         -----
77         x : np array or list
78             array or list of input to the layer.
79
80         Returns
81         -----
82         np array
83             array of output from the layer.
84
85         """
86         x = np.append([1], x)
87         s = x @ self.weights
88         self.last_activation = self._apply_activation(s)
89         return self.last_activation
90
91     def _apply_activation(self, s):
92         """
93         calculate activated output
94
95         Parameters
96         -----
97         s : np array
98             array of the inner product between input and weights.
99
100        Returns
101        -----
102        np array
103            activated output.
104
```

```

105
106
107     if self.activation == 'relu':
108         return np.maximum(s,0)
109     elif self.activation == 'tanh':
110         return np.tanh(s)
111     elif self.activation == 'sigmoid':
112         return 1.0 / (1.0 + np.exp(-s))
113     else:
114         return s # if no or unkown activation, f = s
115
116 def apply_activation_derivative(self, s):
117     """
118     calculate the derivative of activation function
119
120     Parameters
121     -----
122     s : np array
123         array of the inner product between input and weights.
124
125     Returns
126     -----
127     np array
128         calcualted output after activation dirivative.
129
130     """
131     if self.activation == 'relu':
132         grad = np.copy(s)
133         grad[s>0] = 1.0
134         grad[s<=0] = 0.0
135         return grad
136     elif self.activation == 'tanh':
137         return 1 - s ** 2
138     elif self.activation == 'sigmoid':
139         return s * (1-s)
140     else:
141         return np.ones_like(s) # if no or unkown activation, f' = 1
142
143 def update_weights(self, dweights):
144     """
145     Used for gradient descent to change weight values
146
147     Parameters
148     -----
149     dweights : TYPE
150         DESCRIPTION.
151
152     Returns
153     -----
154     None.
155
156     """
157     if self.trainable:
158         self.weights += dweights
159         self.last_dweights = dweights
160
161 def io(self):
162     """
163     Get the input and output number of the layer
164
165     Returns
166     -----
167     int
168         input dimension.
169     TYPE
170         output dimension.
171
172     """
173     return (self.weights.shape[0]-1, self.weights.shape[1])
174
175
176
177
178
179
180
181
182
183
184 class NeuralNetwork(object): # neural network model
185     def __init__(self):
186         """
187         Initialize model with a buffer for layers
188
189     Returns
190     -----
191     None.
192
193     """
194         self._layers = []
195         self.learning_rate = None
196         self.momentum = None
197         self.train_errors = []
198
199     def add_layer(self, layer):
200         """
201         Append a layer at the end
202
203     Parameters
204     -----
205     layer : XXLayer type
206         a nn layer.
207
208     Returns
209     -----
210     None.
211

```

```

212     """
213     self._layers.append(layer)
214
215
216 def _feed_forward(self,x):
217     """
218     Calculate output from an input
219
220     Parameters
221     -----
222     x : np array
223         input vector into the network.
224
225     Returns
226     -----
227     x : np array
228         output vector out of the network.
229
230     """
231     for layer in self._layers:
232         x = layer.call(x)
233     return x
234
235 def _back_prop(self, x, y, learning_rate, momentum=0.0, threshold=0.0, weight_decay=0.0):
236     """
237     Implement back propagation with momentum gradient descent and
238     thresholded output
239
240     Parameters
241     -----
242     x : np array
243         input vector to the network.
244     y : np array
245         target output vector for the network.
246     learning_rate : float
247         learning rate.
248     momentum : float, optional
249         alpha value to control the momentum gradient descent. The default is 0.
250     threshold : float, optional
251         threshold window to be considered 0 or 1, detail see self.train(). The default is 0.
252
253     Returns
254     -----
255     None.
256
257     """
258     output = self._feed_forward(x)
259     # Calculate gradients
260     for i in reversed(range(len(self._layers))): # start from the last layer
261         layer = self._layers[i]
262         if i == len(self._layers) - 1: # for output layer
263             raw_error = y - output
264             raw_error = [0 if np.abs(e)<threshold else e for e in raw_error] # implement thresholding
265             layer.error = raw_error
266             layer.delta = layer.apply_activation_derivative(output) * layer.error
267         else: # for hidden layers
268             next_layer = self._layers[i+1]
269             layer.error = next_layer.weights[1:,:] @ next_layer.delta
270             layer.delta = layer.apply_activation_derivative(layer.last_activation) * layer.error
271     # Update weights
272     for i,layer in enumerate(self._layers):
273         pre_synaptic = (x if i == 0 else self._layers[i-1].last_activation)
274         pre_synaptic = np.append([1], pre_synaptic)
275         pre_synaptic = np.atleast_2d(pre_synaptic)
276         delta_weights = pre_synaptic.T @ np.atleast_2d(layer.delta) * learning_rate # basic gradient descent
277         delta_weights -= 2*weight_decay*learning_rate*layer.weights # implement weight decay
278         delta_weights += momentum * layer.last_dweights # implement momentum
279         layer.update_weights(delta_weights)
280
281
282 def train(self, X_train, Y_train, learning_rate, max_epochs, classify=False,
283          momentum=0, threshold=0, validation_ratio=0.0, weight_decay=0.0,
284          earlystop=None):
285     """
286     Train the network with given input, output, and hyper-parameters
287
288     Parameters
289     -----
290     X_train : list or np array
291         a batch of input vector to the network.
292     Y_train : list or np array
293         a batch of target output for the network.
294     learning_rate : float
295         specifies the learning rate of gradient descent.
296     max_epochs : int
297         specifies the max amount of epochs to train.
298     momentum : float, optional
299         specifies the alpha value for gradient descent. The default is 0.
300     threshold : float, optional
301         specified the threshold windows for the output to consider 0 or 1.
302         Output is 0 if 0<= output < threshold; output is 1 if
303         1-threshold < output <=1.
304         The default is 0.
305     stochastic_ratio : float, optional
306         specifies how much of the input batch is selected.
307         The default is 1.0.
308     earlystop : set of 2 elements, optional
309         specifies earlystop. [0] represents the max value for the output
310         to be the 'same'. [1] represents the patience. The default is None.
311
312     Returns
313     -----
314     errors : np array
315         errors every 10 epochs of training.
316
317     """
318     if earlystop is None:

```

```

319         earllystop = (0, max_epochs//10)
320
321
322     X_train = np.array(X_train)
323     Y_train = np.array(Y_train)
324
325     if X_train.shape[1] != self._layers[0].weights.shape[0]-1:
326         raise ValueError("Input data does not match layer dimension")
327     if Y_train.shape[1] != self._layers[-1].weights.shape[1]:
328         raise ValueError("Output data does not match layer dimension")
329
330     self.learning_rate = learning_rate
331     self.momentum = momentum
332     self.weight_decay = weight_decay
333
334     if classify:
335         validation_i, realtrain_i = stratify_split(Y_train, validation_ratio)
336     else:
337         shuffle_i = get_rand_list(len(X_train))
338         realtrain_i = shuffle_i[int(len(X_train)*validation_ratio):]
339         validation_i = shuffle_i[:int(len(X_train)*validation_ratio)]
340
341     X_vali = [X_train[i] for i in validation_i]
342     Y_vali = [Y_train[i] for i in validation_i]
343     good_layers = self._layers
344     errors = []
345     earllystop_counter = 0
346
347     for epoch in trange(max_epochs+1, ncols=75, unit='epochs'):
348         if epoch % 10 == 0:
349
350             if classify:
351                 error = self.classify_test(X_vali, Y_vali)
352             else:
353                 error = self.raw_test(X_vali, Y_vali)
354             errors.append(error)
355
356
357             if epoch == 0:
358                 print("\nLoss = {} at epoch {}".format(errors[-1], epoch))
359                 continue
360
361             if (np.min(errors[:-1]) - errors[-1]) < earllystop[0]:
362                 earllystop_counter += 1
363                 if earllystop_counter == earllystop[1]:
364                     print("\nEarly stop triggered at epoch {}, restored to epoch {}".format(epoch, epoch-earllystop[1]*10))
365                     self._layers = good_layers
366                     self.train_errors = np.array(errors)
367                     return self.train_errors
368                 else:
369                     good_layers = self._layers
370                     earllystop_counter = 0
371
372             print("\nLoss = {} at epoch {}, training stops in {} epochs".format(errors[-1], epoch, (earllystop[1]-earllystop_counter)*10))
373
374             for i in realtrain_i:
375                 self._back_prop(X_train[i], Y_train[i], learning_rate, momentum, threshold, weight_decay)
376
377
378     self.train_errors = np.array(errors)
379     return np.array(errors)
380
381 def raw_test(self, X_test, Y_test):
382     """
383     Test the model with given data, calculate J2 loss
384
385     Parameters
386     -----
387     X_test : 2D list or np array
388             input data to the network.
389     Y_test : 2D list or np array
390             true output data.
391
392     Returns
393     -----
394     float
395         average J2 loss.
396
397     """
398     X_test = np.array(X_test)
399     Y_test = np.array(Y_test)
400     error = 0
401     for i in range(len(X_test)):
402         pred = self._feed_forward(X_test[i])
403         error += np.sum((pred-Y_test[i])**2)
404     return 0.5*error/len(X_test)
405
406
407 def classify_test(self, X_test, Y_test):
408     """
409     Test the network with given input, output and accuracy
410
411     Parameters
412     -----
413     X_test : list or np array
414             input vector to the network.
415     Y_test : list or np array
416             ground truth for the testing.
417
418     Returns
419     -----
420     float
421         test accuracy
422
423     """

```

```

425     errors = []
426     X_test = np.array(X_test)
427     Y_test = np.array(Y_test)
428     for i in range(len(X_test)):
429         pred = self._feed_forward(X_test[i])
430         pred = np.argmax(pred)
431         truth = np.argmax(Y_test[i])
432         errors.append(pred==truth)
433     return 1-np.sum(errors)/len(errors)
434
435 def get_cm(self, X_test, Y_test):
436     """
437     Give the confusion matrix for classification problems
438
439     Parameters
440     -----
441     X_test : list or np array
442             input vector to the network.
443     Y_test : list or np array
444             ground truth for the testing.
445
446     Returns
447     -----
448     cm : np array
449         confusion matrix.
450
451     """
452     X_test = np.array(X_test)
453     Y_test = np.array(Y_test)
454     n_classes = Y_test.shape[1]
455     cm = np.zeros((n_classes, n_classes))
456     for i in range(len(X_test)):
457         pred = self._feed_forward(X_test[i])
458         max_pred = np.max(pred)
459         pred_bin = np.atleast_2d([p==max_pred for p in pred]).T
460         truth = np.atleast_2d(Y_test[i])
461         cm += pred_bin @ truth
462     return cm
463
464 def save(self, file_name):
465     """
466     Save the model in a json file
467     First line of json is meta data
468     Following line includes layer info
469
470     Parameters
471     -----
472     file_name : str
473         string to save data into.
474
475     Returns
476     -----
477     None.
478
479     """
480     with open(file_name, 'w') as f:
481         meta_dict = {}
482         meta_dict['learning_rate'] = self.learning_rate
483         meta_dict['momentum'] = self.momentum
484         meta_dict['weight_decay'] = self.weight_decay
485         meta_dict['train_errors'] = self.train_errors.tolist()
486         json.dump(meta_dict, f)
487         f.write("\n")
488         for layer in self._layers:
489             layer_dict = {}
490             layer_dict['n_input'], layer_dict['n_neurons'] = layer.io()
491             layer_dict['activation'] = layer.activation
492             layer_dict['trainable'] = layer.trainable
493             layer_dict['weights'] = layer.weights.tolist()
494             json.dump(layer_dict, f)
495             f.write("\n")
496
497 def load(self, file_name):
498     """
499     Loads the json file for a model
500
501     Parameters
502     -----
503     file_name : TYPE
504         DESCRIPTION.
505
506     Returns
507     -----
508     None.
509
510     """
511     with open(file_name, 'r') as f:
512         for i, line in enumerate(f):
513             if i == 0:
514                 meta = json.loads(line)
515                 self.learning_rate = meta['learning_rate']
516                 self.momentum = meta['momentum']
517                 self.weight_decay = meta['weight_decay']
518                 self.train_errors = np.array(meta['train_errors'])
519             else:
520                 layer = json.loads(line)
521                 self.add_layer(DenseLayer(n_input=layer['n_input'],
522                                           n_neurons=layer['n_neurons'],
523                                           activation=layer['activation'],
524                                           weights=layer['weights'],
525                                           trainable=layer['trainable']))
526
527 def info(self):
528     """
529     Print the model information
530
531     Returns
532     -----

```

```

532     None.
533
534     """
535     print('Learning rate: {}\nMomentum: {}'.format(self.learning_rate, self.momentum))
536     for i, layer in enumerate(self._layers):
537         print("Layer {} = input: {}, output: {}, activation: {}, trainability: {}".format(i, layer.io()[0], layer.io()[1],
layer.activation, layer.trainable))
538
539 def layers(self, n):
540     """
541     Get layers of a model
542
543     Parameters
544     -----
545     n : int
546         index of the layer starting from 0.
547
548     Returns
549     -----
550     Layer object
551         the n-th layer of the model.
552
553     """
554     return self._layers[n]
555
556 def predict(self, X):
557     """
558     Make prediction from the given input
559
560     Parameters
561     -----
562     X : 2D list or np array
563         Input data to the network.
564
565     Returns
566     -----
567     pred : list
568         predicted output.
569
570     """
571     pred = []
572     for x in X:
573         for layer in self._layers:
574             x = layer.call(x)
575         pred.append(x)
576     return pred

```

## Appendix C Python Code: p1\_train.py

```
1  #!/usr/bin/env python3
2  # -*- coding: utf-8 -*-
3  """
4  Created on Sun Oct 25 21:27:27 2020
5
6  @author: liu
7  """
8
9  from preprocess import get_train, get_test
10 from settings import SIZES, H3P1_NN, HIDDEN_NEURONS, MAX_EPOCHS, VALI_R
11 from nn import NeuralNetwork, DenseLayer
12
13 train_db = get_train()
14 test_db = get_test()
15 network = NeuralNetwork()
16 network.add_layer(DenseLayer(n_input=SIZES['x'][1], n_neurons=HIDDEN_NEURONS,
17                               activation='sigmoid'))
18 network.add_layer(DenseLayer(n_input=HIDDEN_NEURONS, n_neurons=SIZES['classes'],
19                               activation='sigmoid'))
20 network.train(train_db['x'], train_db['y'], max_epochs=MAX_EPOCHS,
21              classify=True, threshold=0.25,
22              validation_ratio=VALI_R, earlystop=(0,5),
23              learning_rate = 0.04,
24              momentum=0.8,
25              weight_decay=1E-4,
26              )
27 network.save(str(H3P1_NN))
```

## Appendix D Python Code: p1\_test.py

```
1 #!/usr/bin/env python3
2 # -*- coding: utf-8 -*-
3 """
4 Created on Tue Oct 27 18:07:34 2020
5
6 @author: liu
7 """
8 from preprocess import get_train, get_test, find_ES
9 from settings import CLASSES, H3P1_NN, H3P1_CM_PLOT, H3P1_TRAIN_PLOT, MAX_EPOCHS
10 from nn import NeuralNetwork
11 import numpy as np
12 import matplotlib.pyplot as plt
13
14 # Load the network
15 train_db = get_train()
16 test_db = get_test()
17 network = NeuralNetwork()
18 network.load(str(H3P1_NN))
19
20 # Plot training error vs epoch
21 train_errors = network.train_errors
22 epochs = 10*np.arange(len(train_errors))
23 fig1, ax1 = plt.subplots(figsize=(8,6))
24 ax1.plot(epochs, train_errors)
25 ax1.set_xticks(np.append(np.arange(0, epochs[-1], 20), epochs[-1]))
26 ax1.set_title("2-Layer NN Classifier: Error vs Epoch")
27 ax1.set_xlabel("Epoch")
28 ax1.set_ylabel("Train Error")
29 notetxt = find_ES(train_errors, MAX_EPOCHS)
30 fig1.text(0.02, 0.01, "* Note:" + notetxt, ha='left')
31 fig1.tight_layout(rect=[0, 0.02, 1, 1])
32 fig1.savefig(H3P1_TRAIN_PLOT)
33
34 # Plot confusion metrix
35 cm = []
36 cm.append(network.get_cm(train_db['x'], train_db['y']))
37 cm.append(network.get_cm(test_db['x'], test_db['y']))
38 errors = []
39 errors.append(network.classify_test(train_db['x'], train_db['y']))
40 errors.append(network.classify_test(test_db['x'], test_db['y']))
41 fig2, ax2 = plt.subplots(1, 2, figsize=(12, 6))
42 for n in range(2):
43     ax2[n].imshow(cm[n], cmap='Blues')
44     ax2[n].set_xticks(CLASSES)
45     ax2[n].set_yticks(CLASSES)
46     ax2[n].set_xticklabels(CLASSES)
47     ax2[n].set_yticklabels(CLASSES)
48     ax2[n].tick_params(axis='both', which='both', length=0)
49     for i in range(len(CLASSES)):
50         for j in range(len(CLASSES)):
51             c = 'w' if cm[n][i, j] >= 50 else 'k'
52             text = ax2[n].text(j, i, int(cm[n][i, j]), ha="center", va="center", color=c, fontsize=12)
53     ax2[n].set_xlabel("True Class\n({})".format(chr(ord('a')+n)), fontsize=14)
54     ax2[n].set_ylabel("Predicted Class", fontsize=14)
55     for num in CLASSES:
56         ax2[n].axvline(num-0.5, c='cornflowerblue', lw=1.5, alpha=0.3)
57         ax2[n].axhline(num-0.5, c='cornflowerblue', lw=1.5, alpha=0.3)
58 ax2[0].set_title("Result on Train Data (Overall Accuracy = {:.3f})".format(1-errors[0]))
59 ax2[1].set_title("Result on Test Data Result (Overall Accuracy = {:.3f})".format(1-errors[1]))
60 fig2.suptitle("2-Layer NN Classifier: Confusion Matrix")
61 fig2.tight_layout(rect=[0, 0, 1, 0.92])
62
63 fig2.savefig(H3P1_CM_PLOT)
64
65 plt.show()
66 plt.close('all')
```



## Appendix E Python Code: p2\_train.py

```
1 #!/usr/bin/env python3
2 # -*- coding: utf-8 -*-
3
4 from preprocess import get_train, get_test
5 from settings import SIZES, H3P2_NN, HIDDEN_NEURONS, MAX_EPOCHS, VALI_R
6 from nn import NeuralNetwork, DenseLayer
7
8
9 train_db = get_train()
10 test_db = get_test()
11 network = NeuralNetwork()
12 network.add_layer(DenseLayer(n_input=SIZES['x'][1], n_neurons=HIDDEN_NEURONS,
13                               activation='sigmoid'))
14 network.add_layer(DenseLayer(n_input=HIDDEN_NEURONS, n_neurons=784,
15                               activation='sigmoid'))
16 network.train(train_db['x'], train_db['x'], learning_rate = 0.01,
17               max_epochs=MAX_EPOCHS, classify=False, momentum=0.8,
18               threshold=0, validation_ratio=VALI_R, earllystop=(1E-3,5),
19               # weight_decay=1E-5,
20               )
21 network.save(str(H3P2_NN))
```

## Appendix F Python Code: p2\_test.py

```
1  #!/usr/bin/env python3
2  # -*- coding: utf-8 -*-
3  """
4  Created on Tue Oct 27 18:36:30 2020
5
6  @author: liu
7  """
8
9  from preprocess import get_train, get_test, get_rand_list, prepare_img, find_ES
10 from settings import CLASSES, SIZES, MAX_EPOCHS
11 from settings import H3P2_NN, H3P2_TRAIN_PLOT, H3P2_TEST_PLOT, H3P2_FEATURE_MAP, H3P1_NN, H3P1_FEATURE_MAP,
12     HIDDEN_NEURONS
13 from nn import NeuralNetwork
14 import numpy as np
15 import matplotlib.pyplot as plt
16
17 # Load the network
18 train_db = get_train()
19 test_db = get_test()
20 autoenc = NeuralNetwork()
21 autoenc.load(str(H3P2_NN))
22
23 # Plot training error vs epoch
24 train_errors = autoenc.train_errors
25 epochs = 10*np.arange(len(train_errors))
26 fig1, ax1 = plt.subplots(figsize=(8,6))
27 ax1.plot(epochs, train_errors)
28 ax1.set_xticks(np.append(np.arange(0, epochs[-1], 20), epochs[-1]))
29 ax1.set_title("2-Layer NN Autoencoder: Error vs Epoch")
30 ax1.set_xlabel("Epoch")
31 ax1.set_ylabel("Train Error")
32 notetxt = find_ES(train_errors, MAX_EPOCHS)
33 fig1.text(0.02, 0.01, "Note:" + notetxt, ha='left')
34 fig1.tight_layout(rect=[0, 0.02, 1, 1])
35 fig1.savefig(H3P2_TRAIN_PLOT)
36
37 # Plot testing errors
38 test_errors = [[] for _ in CLASSES]
39 train_errors = [[] for _ in CLASSES]
40 for i, x in enumerate(train_db['x']):
41     c = np.argmax(train_db['y'][i])
42     train_errors[c].append(autoenc.raw_test([x], [x]))
43 for i, x in enumerate(test_db['x']):
44     c = np.argmax(test_db['y'][i])
45     test_errors[c].append(autoenc.raw_test([x], [x]))
46 test_errors = np.mean(test_errors, axis=1)
47 test_errors = np.insert(test_errors, 0, np.mean(test_errors))
48 train_errors = np.mean(train_errors, axis=1)
49 train_errors = np.insert(train_errors, 0, np.mean(train_errors))
50 fig2, ax2 = plt.subplots(figsize=(8,6))
51 width = 0.35
52 ticks = [str(c) for c in CLASSES]
53 ticks.insert(0, 'Overall')
54 ax2.bar(np.arange(len(ticks)) - width/2, train_errors, width, label='Train Errors')
55 ax2.bar(np.arange(len(ticks)) + width/2, test_errors, width, label='Test Errors')
56 ax2.set_xticks(np.arange(len(ticks)))
57 ax2.set_xticklabels(ticks)
58 ax2.set_ylabel('Test Error')
59 ax2.set_title("2-Layer NN Autoencoder: Performance on Each Class")
60 ax2.legend(loc='lower right')
61 ax2.grid(axis='y')
62 fig2.tight_layout()
63 fig2.savefig(H3P2_TEST_PLOT)
64
65 # Plot feature maps
66 fig3, ax3 = plt.subplots(5, 4, figsize=(8,10))
67 neuron_i = get_rand_list(HIDDEN_NEURONS)[:20]
68 for i, ni in enumerate(neuron_i):
69     features = autoenc.layers(0).weights[:, ni][1:]
70     features = features / np.linalg.norm(features)
71     ax3[i//4][i%4].imshow(prepare_img(features), cmap='binary')
72     ax3[i//4][i%4].set_xticks([])
73     ax3[i//4][i%4].set_yticks([])
74     ax3[i//4][i%4].set_xlabel('{}{}'.format(chr(ord('a')+i)), fontsize=14)
75 fig3.suptitle("2-Layer NN Autoencoder: Hidden Neuron Features")
76 fig3.tight_layout(rect=[0, 0, 1, 0.95])
77 fig3.savefig(H3P2_FEATURE_MAP)
78
79 classifier = NeuralNetwork()
80 classifier.load(str(H3P1_NN))
81 fig4, ax4 = plt.subplots(5, 4, figsize=(8,10))
82 for i, ni in enumerate(neuron_i):
83     features = classifier.layers(0).weights[:, ni][1:]
84     ax4[i//4][i%4].imshow(prepare_img(features), cmap='binary')
85     ax4[i//4][i%4].set_xticks([])
86     ax4[i//4][i%4].set_yticks([])
87     ax4[i//4][i%4].set_xlabel('{}{}'.format(chr(ord('a')+i)), fontsize=14)
88 fig4.suptitle("2-Layer NN Classifier: Hidden Neuron Features")
89 fig4.tight_layout(rect=[0, 0, 1, 0.95])
90 fig4.savefig(H3P1_FEATURE_MAP)
91
92 # Plot sample output
93 img_i = get_rand_list(SIZES['test'][:8])
94 fig5, ax5 = plt.subplots(2, 8, figsize=(16,4))
95 for i, ii in enumerate(img_i):
96     original = test_db['x'][ii]
97     ax5[0][i].imshow(prepare_img(original), cmap='binary')
98     reconstructed = autoenc.predict([test_db['x'][ii]])
99     ax5[1][i].imshow(prepare_img(reconstructed), cmap='binary')
100    ax5[0][i].set_xticks([])
101    ax5[0][i].set_yticks([])
102    ax5[0][i].set_xlabel('{}{}'.format(chr(ord('a')+i)), fontsize=14)
103    ax5[1][i].set_xticks([])
104    ax5[1][i].set_yticks([])
```

```

104     ax5[1][i].set_xlabel('{}{}'.format(chr(ord('i')+i)), fontsize=14)
105 ax5[0][0].set_ylabel("Original")
106 ax5[1][0].set_ylabel("Reconstructed")
107 fig5.suptitle("2-Layer NN Autoencoder: Sample Outputs")
108 fig5.tight_layout(rect=[0, 0, 1, 0.95])
109 fig5.savefig(H3P2_OUTPUT_MAP)
110
111
112 plt.show()
113 plt.close('all')

```