

EECE6036: Intelligent Systems
Homework # 4

Zuguang Liu (M10291593)

April 23, 2021

1 Denoising Autoencoder

1.1 Problem Statement

An autoencoder shall be implemented by a two-layer neural network to remove noise applied to the MNIST dataset. Then the performance and the model itself is compared with the autoencoder constructed in Homework #3, whose purpose was simply to reconstruct the images.

To tell apart the two models, from now on the report uses **denoising autoencoder** to address this new network, and call the previous one **reconstructing autoencoder**.

1.2 System Description

Using the same stratified sampling policy the previous Homework used, the preprocessing section partitions the total 5,000 data points into 4,000 for training and 1,000 for testing, where there are 400 and 100 images for each model respectively. Over the training, each image is tempered with noise by a certain policy as input, and the original image is learned as target to fulfill the denoising purpose.

The noise policy generates “salt and pepper” noise (salt pixels = 1, pepper pixels = 0) whose densities are approximately 5.56% and 44.44%. This is due to the fact that the dark (< 0.5) vs bright pixels (> 0.5) in an image is at a ratio around 8:1 on average, and the goal of total noise density = 50%.

The autoencoder is implemented by a 2-layer neural network with a hidden layer and an output layer. **The hidden layer incorporates 128 neurons** that transfer the noisy image into a 128-dimensional feature space. Then the output layer brings the reduced feature space back into the original 784 pixel values but with reduced noise. **Both layers use the *sigmoid* activation function (1) and a weight initialization scheme known as “Xavier initialization”** showed in (2), where $\mathcal{U}[a, b]$ is the uniform distribution within the interval between a and b , n_{in} is the input dimension of a layer, and n_{out} is the output dimension of a layer [1].

$$f(x) = \frac{1}{1 + e^{-x}}, \quad f'(x) = f(1 - f) \quad (1)$$

$$W \sim \mathcal{U}\left[-\sqrt{\frac{6}{n_{in} + n_{out}}}, \sqrt{\frac{6}{n_{in} + n_{out}}}\right] \quad (2)$$

The training of the model is done by back propagation per data point, repeated for multiple epochs. 1,000 data points are separated for validation, presented to the model every 10 epochs, resulting in a series of on-line training errors. Within one epoch, the remaining 3,000 points are shuffled, then used to adjust weights and biases. The error is calculated with the average J_2 loss over all data points in the validation set, shown in (3), where N is the number of data points, y_i^n and \hat{y}_i^n are the original and predicted i -th pixel value on the n -th image, respectively.

$$\bar{J}_2 = \frac{1}{N} \sum_{n=1}^N J_2 = \frac{1}{1000} \sum_{n=1}^{1000} \left(\frac{1}{2} \sum_{i=1}^{784} (y_i^n - \hat{y}_i^n)^2 \right) \quad (3)$$

To further improve the training, several mechanisms are used in the training algorithm, including:

- a. Using **Weight decay** for regularization by adding a weight penalty term to the loss function, as in (4) where $\lambda = 10^{-4}$.

$$J = \bar{J}_2 + \lambda \sum_{Layers L} \sum_{j \in Layer L} \sum_{i \in Layer L+1} w_{ij}^2 \quad (4)$$

- b. The gradient decent has an additional term to implement **momentum**, demonstrated in (5), where $\eta = 0.05$, $\alpha = 0.8$.

$$\Delta w_{ij}(t) = -\eta \frac{\partial J}{\partial w_{ij}} + \alpha \Delta w_{ij}(t-1) \quad (5)$$

- c. Though the total repetition is 500 epochs, **an early stopping policy** is used so that it stops when the on-line training loss **does not improve more than 10^{-3} for 50 epochs** compared to the minimum loss over the whole training session. As the validation set used for on-line testing is separate from the data used in back propagation, this ensures the model does not overfit the data.

After training, the resulting model is analyzed by feeding in the test dataset. A reconstructing autoencoder with the same hyper-parameters and policies except zero noise is also trained and tested for comparison.

1.3 Results

For the sake of discussion, both autoencoders are included in plots. Though the prediction is compared against the original image for both models, the reconstructing autoencoder has the exact same image as input, but the denoising autoencoder has noisy images as input.

1.4 Overall Performance

Fig. 1 demonstrates the performance of the autoencoders overall and on each class, where the error is calculated by (3). The on-line training error vs epochs is shown in Fig. 2.




Fig. 1: Performance of the autoencoders on the training and test set

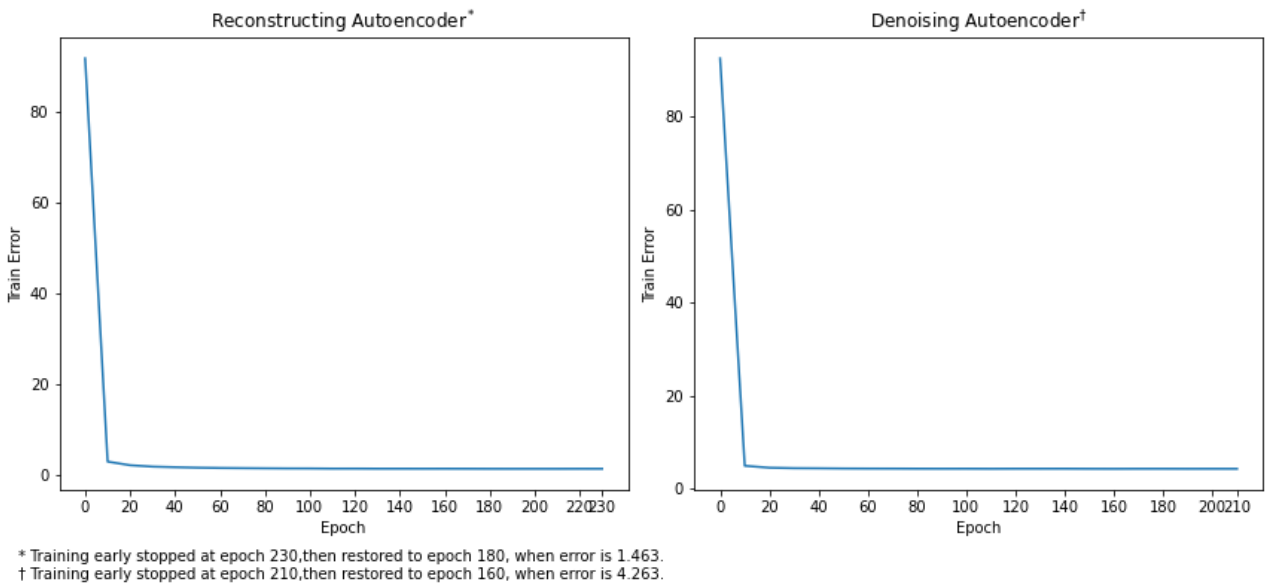


Fig. 2: On-line training error vs training epochs over time

1.5 Features

Weights of 20 neurons in the hidden layer of the two networks are illustrated by 28×28 images in Fig. 3, as the feature space is 784-dimensional, in accordance to 784 pixels of the original images. Though neurons are chosen randomly, the selection of neurons in both model uses the same indexes, so that neurons in the same position are compared.

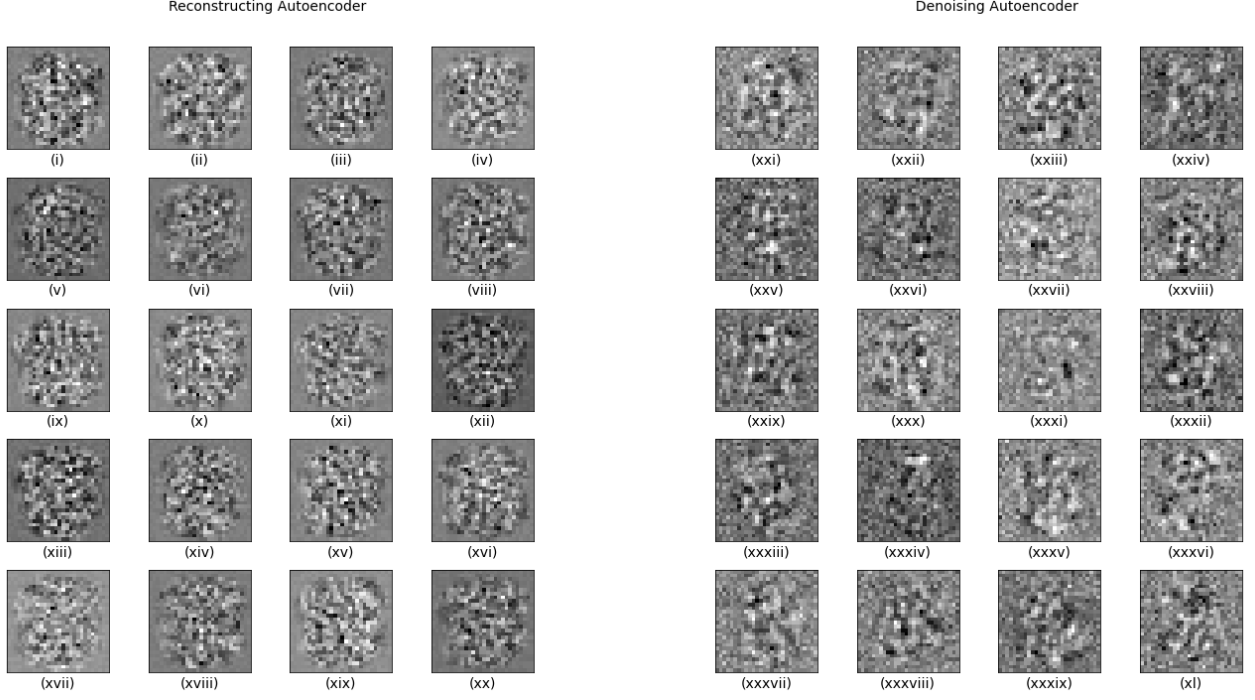


Fig. 3: Feature map of 20 neurons in the hidden layer of both autoencoders.

1.6 Sample Outputs

Fig. 4 uses 8 sets of images to visualize the two model's performance in reconstruction and noise removal.




Fig. 4: Sample output of the models, (ix) - (xvi) for the reconstructing autoencoder, and (ix) - (xvi) for the denoising autoencoder

1.7 Analysis of Results

Observed from Fig. 4, both autoencoders serve their respective purposes well. Comparing the denoised images in with the original images, the “salt and pepper” noise is significantly reduced with minor errors. It is found that images with noise more scattered such as (xxiii) are denoised better than those with clustered noise such as (xxvii), due to the fact that clustered “salt” or “pepper” could be recognized as original pattern to remain.

Comparing the feature maps in their hidden layers, first difference to notice is that there are granular difference in each pixels throughout in the denoising autoencoder, which is absent in the reconstructing autoencoder. Moreover, larger clouds of dark or light pixels are in denoising autoencoder compared to the reconstructing one. It is theorized that the two differences could imply the denoising model learned a “bigger picture” in the image compared to reconstructing model, despite having larger training and testing error shown in Fig. 1 and 2.

2 Classifiers Based on Autoencoders

2.1 Problem Statement

Using the first layer from the autoencoders, two classifiers are constructed with an output layer that predicts the label of images instead. After training, the performance of the two networks shall be compared and contrasted together with the classifier from Homework #3 whose hidden layer is initialized randomly.

To tell apart the models, the report uses **denoising classifier** to address the model based on denoising autoencoder, and calls the reconstruction-based network **reconstructing classifier**. The other one from Homework #3 whose hidden layer is initialized randomly shall be called **BP classifier** as the gradient back-propagates to the first layer to adjust its weights, while only the output layer applies the weight change for the other two models.

2.2 System Description

To control the training settings, **all three classifiers** (reconstructing, denoising and BP) **are trained with exactly the same algorithm and hyper-parameters**.

As the problem states, **the two classifiers use a hidden layer with 128 neurons** from the denoising and reconstructing autoencoder, and an output layer that classifies images with an array of 10 elements, each of which represents the probability that the image can belong to a class. **This output layer is again implemented by sigmoid (1) and Xavier initialization policy (2).**

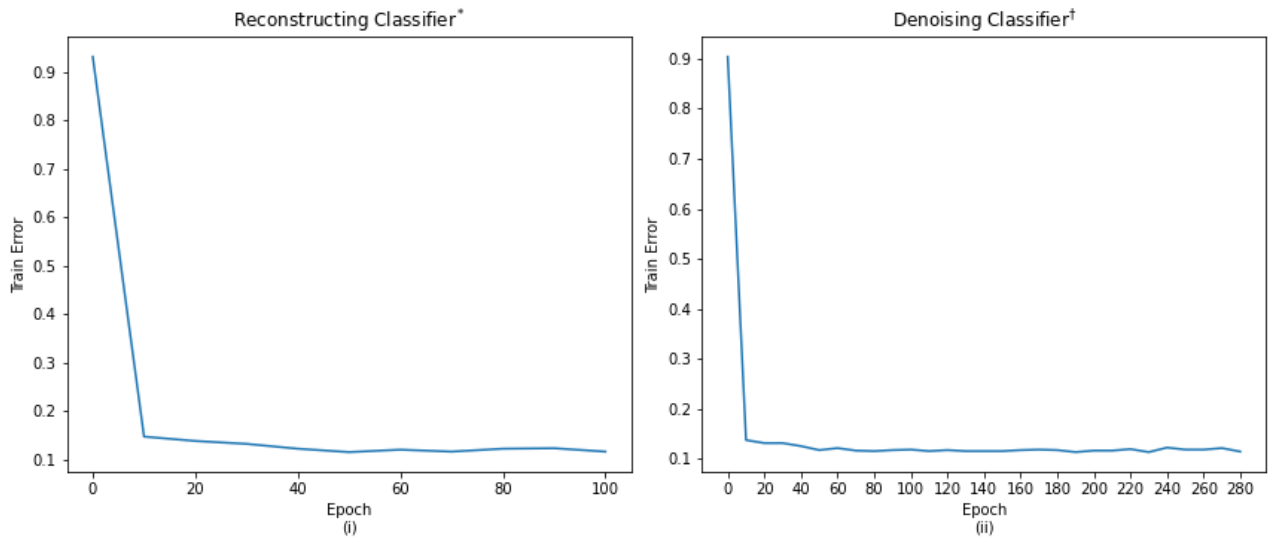
The training of the models proceeds similarly to that of the autoencoders, except the prediction is compared against the labels of the images rather than the image itself. Original version of the images is used as input for both models, since it could be unfair for the reconstructing and BP classifier. There are other minor modifications in the training session, including:

1. The back propagation uses momentum-based gradient descent (5) with $\eta = 0.01$, $\alpha = 0.8$ and weight decay uses $\lambda = 10^{-5}$.
2. Only weights of the output layer change over training. The hidden layer is treated as “read-only”.
3. The loss is calculated by (1 - balanced accuracy), where the balanced accuracy is the hit rate when comparing the true class and the predicted class using “winner-take-all” strategy over the output array.
4. **Operating thresholds of 0.25 and 0.75 are used** so that output $\in [0, 0.25)$ is considered 0 when the corresponding truth is 0, and output $\in (0.75, 1]$ is considered 1 when the corresponding truth is 1.
5. Early stop strategy is again used for regularization. **The training stops when no improvement is observed for 50 epochs.**

After training, test dataset is presented to all models including reconstructing classifier, denoising classifier and BP classifier.

2.3 Results

Time series of on-line training loss on both models is presented in Fig. 5. Fig. 6 shows the confusion matrices of both classifiers on train data and test data, along with the ones of BP classifier.



* Training early stopped at epoch 100, then restored to epoch 50, when error is 0.115.
† Training early stopped at epoch 280, then restored to epoch 230, when error is 0.114.

Fig. 5: Error vs epochs during training of reconstructing classifier, denoising classifier and BP classifier

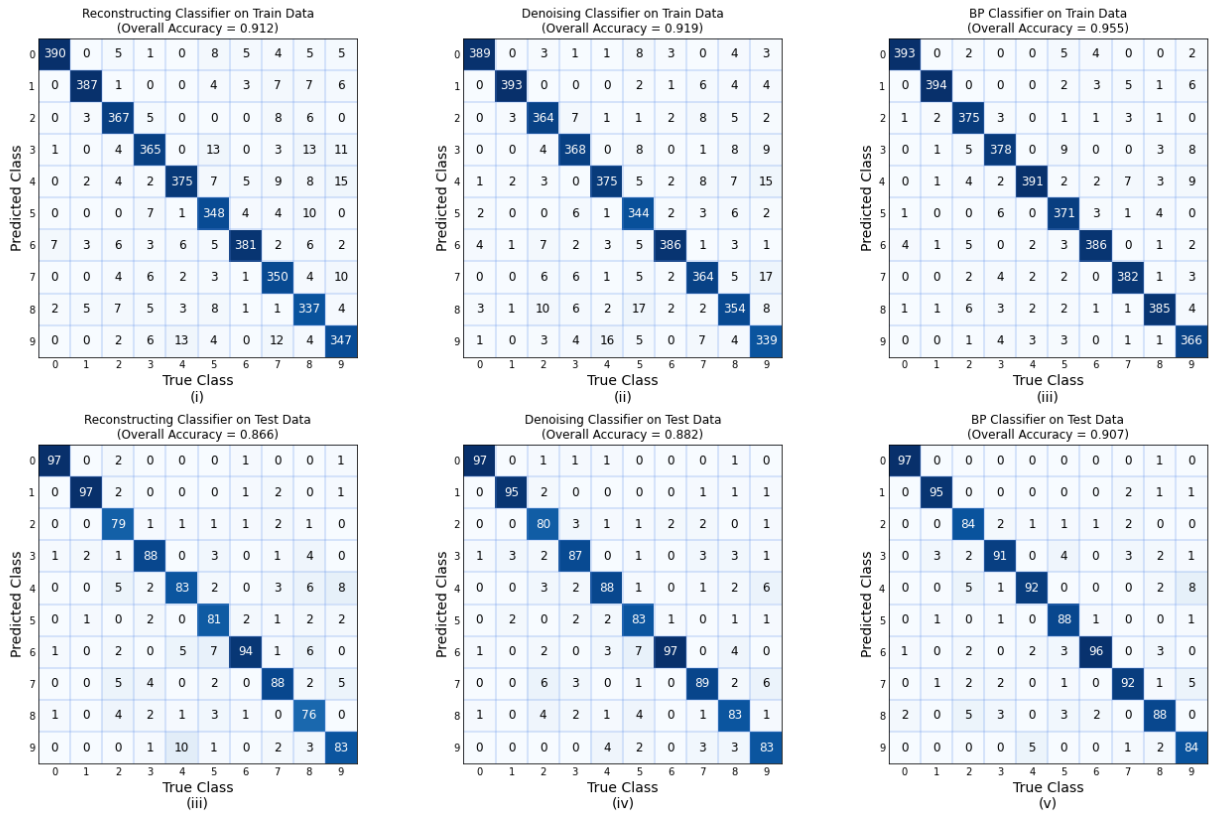


Fig. 6: Performance of reconstructing classifier, denoising classifier and BP classifier

2.4 Analysis of Results

Unfortunately, neither the reconstructing nor the denoising classifier beat the BP classifier in performance (Fig. 6). If to rank their accuracy in classifying, BP is better than denoising, and the reconstructing one is the worst. To break down the analysis, I will discuss the difference between each other separately.

The fact that reconstructing classifier performs worse than the denoising classifier on all classes is not surprising. As the hypothesis in Section 1.4 states, the denoising autoencoder generalizes the image more than the reconstructing autoencoder, and thus, each perceptron in the hidden layer is less likely to be restricted to a small amount of pixels. Comparatively, perceptrons in the reconstructing autoencoder are more likely to miss the global information, which could be important in classification problems.

The denoising classifier has the potential to perform better than the BP classifier, since the train error experienced an amount of up-and-down's in Fig. 5 (ii). Yet it did not manage to win in the end. This could be caused by the fact that it only has 10 neurons to train, while the BP classifier has 138 (128 in layer 1, 10 in layer 2). The pre-trained well-performing autoencoder layer may have established a good basis for the feature reduction, but the inability to fine tune the hidden 128 neurons makes it hard to achieve high-resolution training no matter how long the training lasts. If the pre-trained autoencoder weights are used as a weight initialization policy rather than a “read-only” layer, it is believe that the model could have even better performance and training efficiency than the BP classifier.

References

- [1] X. Glorot and Y. Bengio, “Understanding the difficulty of training deep feedforward neural networks,” p. 8.

Appendix A Python Code: preprocess.py

```
1 import numpy as np
2 import csv
3 import settings
4 import json
5 import pandas as pd
6 import matplotlib.pyplot as plt
7
8 def prepare_img(img_a):
9     img_a = 1 - np.array(img_a).flatten()
10    img_a = img_a / np.linalg.norm(img_a)
11    return np.reshape(img_a, (-1, int(len(img_a)**0.5)), order='F')
12
13 def get_rand_list(length):
14     return np.random.choice(length,length,replace=False).astype(int)
15
16 def prepare_data():
17     x_db = []
18     with open(str(settings.X_FILE)) as csv_file:
19         csv_reader = csv.reader(csv_file, delimiter='\t')
20         for row in csv_reader:
21             x_db.append([float(x) for x in row])
22
23     y_db = []
24     with open(str(settings.Y_FILE)) as csv_file:
25         csv_reader = csv.reader(csv_file, delimiter='\t')
26         for row in csv_reader:
27             y_db.append(int(row[0]))
28
29     print('Distribution of original dataset:',np.bincount(y_db))
30
31     train_i, test_i = stratify_split(y_db,
32                                     settings.SIZES['train']/settings.SIZES['y'])
33
34     train_db = {'x':[x_db[i] for i in train_i],
35                'y':[y_db[i] for i in train_i]}
36
37     test_db = {'x':[x_db[i] for i in test_i],
38               'y':[y_db[i] for i in test_i]}
39
40
41     print('Distribution of train dataset:',np.bincount(train_db['y']))
42     print('Distribution of test dataset:',np.bincount(test_db['y']))
43
44     test_db['y'] = np.eye(settings.SIZES['classes'])[test_db['y']].tolist()
45     train_db['y'] = np.eye(settings.SIZES['classes'])[train_db['y']].tolist()
46
47     with open(str(settings.TRAIN_FILE),'w') as f:
48         json.dump(train_db, f)
49
50     print("Saved train data in", settings.TRAIN_FILE)
51
52     with open(str(settings.TEST_FILE),'w') as f:
53         json.dump(test_db, f)
54
55     print("Saved test data in", settings.TEST_FILE)
56
57 def stratify_split(y, ratio):
58     if len(np.array(y).shape) > 1: # collapse for one hot
59         y = np.argmax(y,axis=1)
60     df = pd.DataFrame(y).groupby(0) # Sort data by class
61     indxs = [] # buffer for indexes
62     for _,g in df:
63         indxs.append(g.index.to_numpy()) # indexes of each class take a row
64     indxs = np.array(indxs)
65     p1_indx = indxs[:, :int(indxs.shape[1]*ratio)].flatten() # partition 1
66     np.random.shuffle(p1_indx) # mix index
67     p2_indx = indxs[:, int(indxs.shape[1]*ratio):].flatten() # partition 2
68     np.random.shuffle(p2_indx) # mix index
69     return p1_indx, p2_indx
70
71
72 def get_test():
73     with open(str(settings.TEST_FILE),'r') as f:
74         return json.load(f)
75
76 def get_train():
77     with open(str(settings.TRAIN_FILE),'r') as f:
78         return json.load(f)
79
80 def add_noise(img, noise_type):
81     img = np.array(img)
82     noise_type = noise_type.lower()
83     if noise_type == "gaussian":
84         mu = 0.5
85         sigma = np.sqrt(0.001)
86         return img + np.random.normal(mu,sigma, len(img))
87     elif noise_type == "s&p":
88         density = 0.5
89         svp = 1/8
90         rand_i = get_rand_list(len(img))[:int(len(img)*density)]
91         salt_i = rand_i[:int(len(rand_i)*svp/(svp+1.0))]
92         pepper_i = rand_i[int(len(rand_i)/(svp+1.0)):]
93         img[salt_i] = 0.0
94         img[pepper_i] = 1.0
95         return img
96     elif noise_type == "poisson":
97         vals = len(np.unique(img))
98         vals = 2 ** np.ceil(np.log2(vals))
99         return np.random.poisson(img * vals) / float(vals)
100    elif noise_type == "speckle":
101        return img + img*np.random.uniform(0,1)
102
103 def int_to_roman(num):
```

```

105     result = ''
106     mapping = {1000:'M', 900:'CM', 500:'D', 400:'CD', 100:'C', 90:'XC', 50:'L', 40:'XL', 10:'X', 9:'IX', 5:'V', 4:'IV', 1:'I'}
107
108     while num != 0:
109         for k, v in mapping.items():
110             if num >= k:
111                 dividend = int(num/k)
112                 num %= k
113                 result += dividend*v
114     return result.lower()
115
116 if __name__ == '__main__':
117     # prepare_data()
118
119     x = get_train()['x']
120     i = int(get_rand_list(len(x))[0])
121     fig, ax = plt.subplots(1,2,figsize=(8,6))
122     ax[0].imshow(prepare_img(x[i],cmap='binary')
123     ax[1].imshow(prepare_img(add_noise(x[i],"s&p")), cmap='binary')
124
125     # print(int_to_roman(15))
126     pass

```

Appendix B Python Code: nn.py

```
1 import numpy as np
2 from preprocess import get_rand_list, stratify_split
3 import json
4 from tqdm import trange
5
6
7 class DenseLayer(object): # fully connected layer
8     def __init__(self, n_input, n_neurons, activation=None, trainable=True,
9                 weights=None):
10         """
11         Initialize a fully-connected layer
12
13         Parameters
14         -----
15         n_input : uint
16             number of input nodes.
17         n_neurons : uint
18             number of neurons / output nodes.
19         activation : str, optional
20             activation function name. The default is None.
21         weights : np array, optional
22             matrix for weights. The default is None.
23
24         Returns
25         -----
26         None.
27
28         """
29
30         if weights is None:
31             a = np.sqrt(6/(n_input+n_neurons))
32             self.weights = np.random.uniform(low=-a, high=+a, size=(n_input+1, n_neurons)) #Xavier initialization
33         else:
34             weights = np.array(weights)
35             if weights.shape == (n_input+1, n_neurons):
36                 self.weights = weights
37             else:
38                 raise ValueError("Given weights does not match given dimensions")
39         self.trainable = trainable
40
41         self.last_dweights = np.zeros((n_input+1, n_neurons))
42
43         self.activation = activation
44         self.last_activation = None
45         self.error = None
46         self.delta = None
47
48     def set_trainable(self, trainable):
49         """
50         Configure if the layer is trainable
51
52         Parameters
53         -----
54         trainable : bool
55             whether the layer is trainable.
56
57         Returns
58         -----
59         None.
60
61         """
62         self.trainable = trainable
63
64     def call(self, x):
65         """
66         Calculate the output given input
67
68         Parameters
69         -----
70         x : np array or list
71             array or list of input to the layer.
72
73         Returns
74         -----
75         np array
76             array of output from the layer.
77
78         """
79         x = np.append([1], x)
80         s = x @ self.weights
81         self.last_activation = self._apply_activation(s)
82         return self.last_activation
83
84     def _apply_activation(self, s):
85         """
86         calculate activated output
87
88         Parameters
89         -----
90         s : np array
91             array of the inner product between input and weights.
92
93         Returns
94         -----
95         np array
96             activated output.
97
98         """
99
100         if self.activation == 'relu':
101             return np.maximum(s, 0)
102         elif self.activation == 'tanh':
103             return np.tanh(s)
104         elif self.activation == 'sigmoid':
```

```

105         return 1.0 / (1.0 + np.exp(-s))
106     else:
107         return s # if no or unknown activation, f = s
108
109 def apply_activation_derivative(self, s):
110     """
111     calculate the derivative of activation function
112
113     Parameters
114     -----
115     s : np array
116         array of the inner product between input and weights.
117
118     Returns
119     -----
120     np array
121         calculated output after activation derivative.
122
123     """
124     if self.activation == 'relu':
125         grad = np.copy(s)
126         grad[s>0] = 1.0
127         grad[s<=0] = 0.0
128         return grad
129     elif self.activation == 'tanh':
130         return 1 - s ** 2
131     elif self.activation == 'sigmoid':
132         return s * (1-s)
133     else:
134         return np.ones_like(s) # if no or unknown activation, f' = 1
135
136 def update_weights(self, dweights):
137     """
138     Used for gradient descent to change weight values
139
140     Parameters
141     -----
142     dweights : TYPE
143         DESCRIPTION.
144
145     Returns
146     -----
147     None.
148
149     """
150     if self.trainable:
151         self.weights += dweights
152         self.last_dweights = dweights
153
154 def io(self):
155     """
156     Get the input and output number of the layer
157
158     Returns
159     -----
160     int
161         input dimension.
162     TYPE
163         output dimension.
164
165     """
166     return (self.weights.shape[0]-1, self.weights.shape[1])
167
168
169
170
171
172
173
174
175
176
177 class NeuralNetwork(object): # neural network model
178     def __init__(self):
179         """
180         Initialize model with a buffer for layers
181
182     Returns
183     -----
184     None.
185
186     """
187     self._layers = []
188     self.learning_rate = None
189     self.momentum = None
190     self.train_errors = []
191
192     def add_layer(self, layer):
193         """
194         Append a layer at the end
195
196     Parameters
197     -----
198     layer : XXLayer type
199         a nn layer.
200
201     Returns
202     -----
203     None.
204
205     """
206     self._layers.append(layer)
207
208
209     def _feed_forward(self, x):
210         """
211         Calculate output from an input

```

```

212
213
214
215     Parameters
216     -----
217     x : np array
218         input vector into the network.
219
220     Returns
221     -----
222     x : np array
223         output vector out of the network.
224
225     """
226     for layer in self._layers:
227         x = layer.call(x)
228     return x
229
230 def _back_prop(self, x, y, learning_rate, momentum=0.0, threshold=0.0, weight_decay=0.0):
231     """
232     Implement back propagation with momentum gradient descent and
233     thresholded output
234
235     Parameters
236     -----
237     x : np array
238         input vector to the network.
239     y : np array
240         target output vector for the network.
241     learning_rate : float
242         learning rate.
243     momentum : float, optional
244         alpha value to control the momentum gradient descent. The default is 0.
245     threshold : float, optional
246         threshold window to be considered 0 or 1, detail see self.train(). The default is 0.
247
248     Returns
249     -----
250     None.
251
252     """
253     output = self._feed_forward(x)
254     # Calculate gradients
255     for i in reversed(range(len(self._layers))): # start from the last layer
256         layer = self._layers[i]
257         if i == len(self._layers) - 1: # for output layer
258             raw_error = y - output
259             raw_error = [0 if np.abs(e) < threshold else e for e in raw_error] # implement thresholding
260             layer.error = raw_error
261             layer.delta = layer.apply_activation_derivative(output) * layer.error
262         else: # for hidden layers
263             next_layer = self._layers[i+1]
264             layer.error = next_layer.weights[i:, :] @ next_layer.delta
265             layer.delta = layer.apply_activation_derivative(layer.last_activation) * layer.error
266     # Update weights
267     for i, layer in enumerate(self._layers):
268         pre_synaptic = (x if i == 0 else self._layers[i-1].last_activation)
269         pre_synaptic = np.append([i], pre_synaptic)
270         pre_synaptic = np.atleast_2d(pre_synaptic)
271         delta_weights = pre_synaptic.T @ np.atleast_2d(layer.delta) * learning_rate # basic gradient descent
272         delta_weights -= 2*weight_decay*learning_rate*layer.weights # implement weight decay
273         delta_weights += momentum * layer.last_dweights # implement momentum
274         layer.update_weights(delta_weights)
275
276 def train(self, X_train, Y_train, learning_rate, max_epochs, classify=False,
277           momentum=0, threshold=0, validation_ratio=0.0, weight_decay=0.0,
278           earllystop=None):
279     """
280     Train the network with given input, output, and hyper-parameters
281
282     Parameters
283     -----
284     X_train : list or np array
285         a batch of input vector to the network.
286     Y_train : list or np array
287         a batch of target output for the network.
288     learning_rate : float
289         specifies the learning rate of gradient descent.
290     max_epochs : int
291         specifies the max amount of epochs to train.
292     momentum : float, optional
293         specifies the alpha value for gradient descent. The default is 0.
294     threshold : float, optional
295         specified the threshold windows for the output to consider 0 or 1.
296         Output is 0 if 0 <= output < threshold; output is 1 if
297         1-threshold < output <=1.
298         The default is 0.
299     stochastic_ratio : float, optional
300         specifies how much of the input batch is selected.
301         The default is 1.0.
302     earllystop : set of 2 elements, optional
303         specifies earllystop. [0] represents the max value for the output
304         to be the 'same'. [1] represents the patience. The default is None.
305
306     Returns
307     -----
308     errors : np array
309         errors every 10 epochs of training.
310
311     """
312     if earllystop is None:
313         earllystop = (0, max_epochs//10)
314
315     X_train = np.array(X_train)
316     Y_train = np.array(Y_train)
317
318     if X_train.shape[1] != self._layers[0].weights.shape[0]-1:

```

```

319         raise ValueError("Input data does not match layer dimension")
320     if Y_train.shape[1] != self._layers[-1].weights.shape[1]:
321         raise ValueError("Output data does not match layer dimension")
322
323     self.learning_rate = learning_rate
324     self.momentum = momentum
325     self.weight_decay = weight_decay
326
327     if classify:
328         validation_i, realtrain_i = stratify_split(Y_train, validation_ratio)
329     else:
330         shuffle_i = get_rand_list(len(X_train))
331         realtrain_i = shuffle_i[int(len(X_train)*validation_ratio):]
332         validation_i = shuffle_i[:int(len(X_train)*validation_ratio)]
333
334     X_vali = [X_train[i] for i in validation_i]
335     Y_vali = [Y_train[i] for i in validation_i]
336     good_layers = self._layers
337     errors = []
338     earllystop_counter = 0
339
340     self.info()
341
342     for epoch in trange(max_epochs+1, ncols=75, unit='epoch'):
343
344         if epoch % 10 == 0:
345
346             if classify:
347                 error = self.classify_test(X_vali, Y_vali)
348             else:
349                 error = self.raw_test(X_vali, Y_vali)
350             errors.append(error)
351
352
353             if epoch == 0:
354                 print("\nLoss = {} at epoch {}".format(errors[-1], epoch))
355                 continue
356
357             if (np.min(errors[:-1]) - errors[-1]) < earllystop[0]:
358                 earllystop_counter += 1
359                 if earllystop_counter == earllystop[1]:
360                     print("\nEarly stop triggered at epoch {}, restored to epoch {}".format(epoch, epoch-earllystop[1]*10))
361                     self._layers = good_layers
362                     self.train_errors = np.array(errors)
363                     return self.train_errors
364                 else:
365                     good_layers = self._layers
366                     earllystop_counter = 0
367
368             print("\nLoss = {} at epoch {}, training stops in {} epochs".format(errors[-1], epoch, (earllystop[1]-earllystop_counter)*10))
369
370     np.random.shuffle(realtrain_i)
371     for i in realtrain_i:
372         self._back_prop(X_train[i], Y_train[i], learning_rate, momentum, threshold, weight_decay)
373
374
375     self.train_errors = np.array(errors)
376     return np.array(errors)
377
378 def raw_test(self, X_test, Y_test):
379     """
380     Test the model with given data, calculate J2 loss
381
382     Parameters
383     -----
384     X_test : 2D list or np array
385             input data to the network.
386     Y_test : 2D list or np array
387             true output data.
388
389     Returns
390     -----
391     float
392         average J2 loss.
393
394     """
395     X_test = np.array(X_test)
396     Y_test = np.array(Y_test)
397     error = 0
398     for i in range(len(X_test)):
399         pred = self._feed_forward(X_test[i])
400         error += np.sum((pred-Y_test[i])**2)
401     return 0.5*error/len(X_test)
402
403
404 def classify_test(self, X_test, Y_test):
405     """
406     Test the network with given input, output and accuracy
407
408     Parameters
409     -----
410     X_test : list or np array
411             input vector to the network.
412     Y_test : list or np array
413             ground truth for the testing.
414
415     Returns
416     -----
417     float
418         test accuracy
419
420     """
421     errors = []
422     X_test = np.array(X_test)
423     Y_test = np.array(Y_test)

```

```

425     for i in range(len(X_test)):
426         pred = self._feed_forward(X_test[i])
427         pred = np.argmax(pred)
428         truth = np.argmax(Y_test[i])
429         errors.append(pred==truth)
430     return 1-np.sum(errors)/len(errors)
431
432 def get_cm(self, X_test, Y_test):
433     """
434     Give the confusion matrix for classification problems
435
436     Parameters
437     -----
438     X_test : list or np array
439             input vector to the network.
440     Y_test : list or np array
441             ground truth for the testing.
442
443     Returns
444     -----
445     cm : np array
446         confusion matrix.
447
448     """
449     X_test = np.array(X_test)
450     Y_test = np.array(Y_test)
451     n_classes = Y_test.shape[1]
452     cm = np.zeros((n_classes, n_classes))
453     for i in range(len(X_test)):
454         pred = self._feed_forward(X_test[i])
455         max_pred = np.max(pred)
456         pred_bin = np.atleast_2d([p==max_pred for p in pred]).T
457         truth = np.atleast_2d(Y_test[i])
458         cm += pred_bin @ truth
459     return cm
460
461 def save(self, file_name):
462     """
463     Save the model in a json file
464     First line of json is meta data
465     Following line includes layer info
466
467     Parameters
468     -----
469     file_name : str
470             string to save data into.
471
472     Returns
473     -----
474     None.
475
476     """
477     if type(file_name) is not str:
478         file_name = str(file_name)
479
480     with open(file_name, 'w') as f:
481         meta_dict = {}
482         meta_dict['learning_rate'] = self.learning_rate
483         meta_dict['momentum'] = self.momentum
484         meta_dict['weight_decay'] = self.weight_decay
485         meta_dict['train_errors'] = self.train_errors.tolist()
486         json.dump(meta_dict, f)
487         f.write("\n")
488         for layer in self._layers:
489             layer_dict = {}
490             layer_dict['n_input'], layer_dict['n_neurons'] = layer.io()
491             layer_dict['activation'] = layer.activation
492             layer_dict['trainable'] = layer.trainable
493             layer_dict['weights'] = layer.weights.tolist()
494             json.dump(layer_dict, f)
495             f.write("\n")
496
497 def load(self, file_name):
498     """
499     Loads the json file for a model
500
501     Parameters
502     -----
503     file_name : TYPE
504             DESCRIPTION.
505
506     Returns
507     -----
508     None.
509
510     """
511     if type(file_name) is not str:
512         file_name = str(file_name)
513
514     with open(file_name, 'r') as f:
515         for i, line in enumerate(f):
516             if i == 0:
517                 meta = json.loads(line)
518                 self.learning_rate = meta['learning_rate']
519                 self.momentum = meta['momentum']
520                 self.weight_decay = meta['weight_decay']
521                 self.train_errors = np.array(meta['train_errors'])
522             else:
523                 layer = json.loads(line)
524                 self.add_layer(DenseLayer(n_input=layer['n_input'],
525                                           n_neurons=layer['n_neurons'],
526                                           activation=layer['activation'],
527                                           weights=layer['weights'],
528                                           trainable=layer['trainable']))
529
530 def info(self):
531     """
532     Print the model information

```

```

532     Returns
533     -----
534     None.
535
536     """
537     print("{} layer neural network".format(len(self._layers)))
538     print('Learning rate: {}\nMomentum: {}\nWeight Decay: {}'.format(self.learning_rate, self.momentum, self.weight_decay)
539 )
540     for i, layer in enumerate(self._layers):
541         print("Layer {} = input: {}, output: {}, activation: {}, trainable: {}".format(i, layer.io()[0], layer.io()[1],
542         layer.activation, layer.trainable))
543
544 def layers(self, n=None):
545     """
546     Get layers of a model
547
548     Parameters
549     -----
550     n : int
551         index of the layer starting from 0.
552
553     Returns
554     -----
555     Layer object
556         the n-th layer of the model.
557
558     """
559     if n is None:
560         return self._layers
561     else:
562         return self._layers[n]
563
564 def predict(self, X):
565     """
566     Make prediction from the given input
567
568     Parameters
569     -----
570     X : 2D list or np array
571         Input data to the network.
572
573     Returns
574     -----
575     pred : list
576         predicted output.
577
578     """
579     pred = []
580     for x in X:
581         for layer in self._layers:
582             x = layer.call(x)
583         pred.append(x)
584     return pred
585
586 def pop_layer(self):
587     self._layers.pop()

```


Appendix C Python Code: p1_train.py

```
1 from preprocess import get_train, get_test, add_noise
2 from settings import SIZES, H4P1_NN, HIDDEN_NEURONS, MAX_EPOCHS, VALI_R, NOISE, PATIENCE
3 from nn import NeuralNetwork, DenseLayer
4
5 train_db = get_train()
6 test_db = get_test()
7 autoenc = NeuralNetwork()
8 autoenc.add_layer(DenseLayer(n_input=SIZES['x'][1], n_neurons=HIDDEN_NEURONS,
9                               activation='sigmoid'))
10 autoenc.add_layer(DenseLayer(n_input=HIDDEN_NEURONS, n_neurons=SIZES['x'][1],
11                              activation='sigmoid'))
12 noisy_x = [add_noise(x, NOISE) for x in train_db['x']]
13 autoenc.train(noisy_x, train_db['x'], max_epochs=MAX_EPOCHS,
14               classify=False,
15               validation_ratio=VALI_R, earlystop=(1E-3, PATIENCE),
16               learning_rate = 0.01,
17               momentum=0.8,
18               weight_decay=1E-4,
19               )
20 autoenc.save(H4P1_NN)
```

Appendix D Python Code: p1_test.py

```
1 from preprocess import get_train, get_test, get_rand_list, prepare_img, add_noise, int_to_roman
2 from settings import CLASSES, SIZES, PATIENCE, NOISE
3 from settings import H4P1_NN, H4P1_TRAIN_PLOT, H4P1_TEST_PLOT, H4P1_FEATURE_MAP, H4P1_OUTPUT_MAP, H3P2_NN, HIDDEN_NEURONS
4 from nn import NeuralNetwork
5 import numpy as np
6 import matplotlib.pyplot as plt
7
8 # Load the network
9 train_db = get_train()
10 test_db = get_test()
11 autoenc_noise = NeuralNetwork()
12 autoenc_noise.load(H4P1_NN)
13 autoenc_clean = NeuralNetwork()
14 autoenc_clean.load(H3P2_NN)
15
16 # Plot training error vs epoch
17 train_errors = autoenc_clean.train_errors, autoenc_noise.train_errors
18 epochs = 10*np.arange(len(train_errors[0])), 10*np.arange(len(train_errors[1]))
19 fig1, ax1 = plt.subplots(1,2,figsize=(12,6))
20 for i in range(2):
21     ax1[i].plot(epochs[i], train_errors[i])
22     ax1[i].set_xticks(np.append(np.arange(0,epochs[i][-1],20),epochs[i][-1]))
23     ax1[i].set_xlabel("Epoch")
24     ax1[i].set_ylabel("Train Error")
25 ax1[0].set_title("Reconstructing Autoencoder$~*$")
26 ax1[1].set_title("Denoising Autoencoder$~ $ ")
27 fig1.text(0.02, 0.01, '* Training early stopped at epoch {},\n'
28         'then restored to epoch {}, when error is {:.3f}.\n'
29         'Training early stopped at epoch {},\n'
30         'then restored to epoch {}, when error is {:.3f}.\n'
31         .format(epochs[0][-1], epochs[0][-1-PATIENCE], train_errors[0][-1-PATIENCE],
32               epochs[1][-1], epochs[1][-1-PATIENCE], train_errors[1][-1-PATIENCE],
33               ), ha='left')
34 fig1.tight_layout(rect=[0, 0.08, 1, 0.95])
35 fig1.savefig(H4P1_TRAIN_PLOT)
36
37 # Plot training errors
38 fig2, ax2 = plt.subplots(1,2, figsize=(16,6))
39
40 for n in range(2):
41     if n == 0:
42         autoenc = autoenc_clean
43         ax2[n].set_title("Reconstructing Autoencoder")
44     else:
45         autoenc = autoenc_noise
46         ax2[n].set_title("Denoising Autoencoder")
47     test_errors = [[] for _ in CLASSES]
48     train_errors = [[] for _ in CLASSES]
49     for i,x in enumerate(train_db['x']):
50         c = np.argmax(train_db['y'][i])
51         if n == 0:
52             train_errors[c].append(autoenc.raw_test([x],[x]))
53         else:
54             train_errors[c].append(autoenc.raw_test([add_noise(x,NOISE)],[x]))
55     for i,x in enumerate(test_db['x']):
56         c = np.argmax(test_db['y'][i])
57         if n==0:
58             test_errors[c].append(autoenc.raw_test([x],[x]))
59         else:
60             test_errors[c].append(autoenc.raw_test([add_noise(x,NOISE)],[x]))
61     test_errors = np.mean(test_errors,axis=1)
62     test_errors = np.insert(test_errors, 0, np.mean(test_errors))
63     train_errors = np.mean(train_errors,axis=1)
64     train_errors = np.insert(train_errors, 0, np.mean(train_errors))
65     width = 0.35
66     ticks = [str(c) for c in CLASSES]
67     ticks.insert(0,'Overall')
68     ax2[n].bar(np.arange(len(ticks)) - width/2, train_errors, width, label='Train Errors')
69     ax2[n].bar(np.arange(len(ticks)) + width/2, test_errors, width, label='Test Errors')
70     ax2[n].set_xticks(np.arange(len(ticks)))
71     ax2[n].set_xticklabels(ticks)
72     ax2[n].set_ylabel('Test Error')
73     ax2[n].set_xlabel('(' + int_to_roman(n+1) + ')')
74     ax2[n].legend(loc='lower right')
75     ax2[n].grid(axis='y')
76
77 fig2.tight_layout(rect=[0, 0, 1, 0.95])
78 fig2.savefig(H4P1_TEST_PLOT)
79
80 # Plot feature maps
81 fig3, ax3 = plt.subplots(5,9, figsize=(18,10))
82 neuron_i = get_rand_list(HIDDEN_NEURONS)[:20]
83 features = [0]*2
84 for i,ni in enumerate(neuron_i):
85     features[0] = autoenc_clean.layers(0).weights[:,ni][1:]
86     features[1] = autoenc_noise.layers(0).weights[:,ni][1:]
87     ax3[i//4][4].axis('off')
88     for j in range(2):
89         ax3[i//4][i%4+5*j].imshow(prepare_img(features[j]), cmap='binary')
90         ax3[i//4][i%4+5*j].set_xticks([])
91         ax3[i//4][i%4+5*j].set_yticks([])
92         ax3[i//4][i%4+5*j].set_xlabel('(' + int_to_roman(i+20*j+1) + ')', fontsize=14)
93
94 fig3.suptitle('Reconstructing Autoencoder{}Denoising Autoencoder'.format(' '*123), fontsize=14)
95 fig3.tight_layout(rect=[0, 0, 1, 0.93])
96 fig3.savefig(H4P1_FEATURE_MAP)
97
98 # Plot sample output
99
100 img_i = get_rand_list(SIZES['test'][:8])
101 fig4, ax4 = plt.subplots(4,8, figsize=(16,8))
102 for i, ii in enumerate(img_i):
103     clean = test_db['x'][ii]
104     ax4[0][i].imshow(prepare_img(clean), cmap='binary')
```

```

105 ax4[0][i].set_xticks([])
106 ax4[0][i].set_yticks([])
107 ax4[0][i].set_xlabel('(' + int_to_roman(i+0*8+1) + ')', fontsize=14)
108
109 reconstructed = autoenc_clean.predict([clean])
110 ax4[1][i].imshow(prepare_img(reconstructed), cmap='binary')
111 ax4[1][i].set_xticks([])
112 ax4[1][i].set_yticks([])
113 ax4[1][i].set_xlabel('(' + int_to_roman(i+1*8+1) + ')', fontsize=14)
114
115 noisy = add_noise(clean, NOISE)
116 ax4[2][i].imshow(prepare_img(noisy), cmap='binary')
117 ax4[2][i].set_xticks([])
118 ax4[2][i].set_yticks([])
119 ax4[2][i].set_xlabel('(' + int_to_roman(i+2*8+1) + ')', fontsize=14)
120
121 denoised = autoenc_noise.predict([noisy])
122 ax4[3][i].imshow(prepare_img(denoised), cmap='binary')
123 ax4[3][i].set_xticks([])
124 ax4[3][i].set_yticks([])
125 ax4[3][i].set_xlabel('(' + int_to_roman(i+3*8+1) + ')', fontsize=14)
126
127 ax4[0][0].set_ylabel("Original", fontsize=14)
128 ax4[1][0].set_ylabel("Reconstructed", fontsize=14)
129 ax4[2][0].set_ylabel("Noisy", fontsize=14)
130 ax4[3][0].set_ylabel("Denoised", fontsize=14)
131 fig4.tight_layout(rect=[0, 0, 1, 0.95])
132 fig4.savefig(H4P1_OUTPUT_MAP)
133
134 plt.show()
135 plt.close('all')

```

Appendix E Python Code: p2_train.py

```
1 from preprocess import get_train
2 from settings import SIZES, H3P2_NN, HIDDEN_NEURONS, MAX_EPOCHS, VALI_R, H4P1_NN, H4P2C1_NN, H4P2C2_NN, PATIENCE
3 from nn import NeuralNetwork, DenseLayer
4
5
6 train_db = get_train()
7
8 nn_clean = NeuralNetwork()
9 nn_clean.load(H3P2_NN)
10 nn_clean.pop_layer()
11 for layer in nn_clean.layers():
12     layer.set_trainable(False)
13 nn_clean.add_layer(DenseLayer(n_input=HIDDEN_NEURONS, n_neurons=SIZES['classes'],
14                               activation='sigmoid'))
15
16 nn_noise = NeuralNetwork()
17 nn_noise.load(H4P1_NN)
18 nn_noise.pop_layer()
19 for layer in nn_noise.layers():
20     layer.set_trainable(False)
21 nn_noise.add_layer(DenseLayer(n_input=HIDDEN_NEURONS, n_neurons=SIZES['classes'],
22                               activation='sigmoid'))
23
24
25 nn_clean.train(train_db['x'], train_db['y'], max_epochs=MAX_EPOCHS,
26               classify=True, threshold=0.25,
27               validation_ratio=VALI_R, earllystop=(0,PATIENCE),
28               learning_rate = 0.002,
29               momentum=0.8,
30               weight_decay=1E-4,
31               )
32 nn_clean.save(H4P2C1_NN)
33
34
35 nn_noise.train(train_db['x'], train_db['y'], max_epochs=MAX_EPOCHS,
36               classify=True, threshold=0.25,
37               validation_ratio=VALI_R, earllystop=(0,PATIENCE),
38               learning_rate = 0.002,
39               momentum=0.8,
40               weight_decay=1E-4,
41               )
42 nn_noise.save(H4P2C2_NN)
```

Appendix F Python Code: p2_test.py

```
1 from preprocess import get_train, get_test, int_to_roman
2 from settings import CLASSES, H4P2C1_NN, H4P2C2_NN, H4P2_CM_PLOT, H4P2_TRAIN_PLOT, PATIENCE, H3P1_NN
3 from nn import NeuralNetwork
4 import numpy as np
5 import matplotlib.pyplot as plt
6
7 # Load the autoenc_clean
8 train_db = get_train()
9 test_db = get_test()
10 autoenc_clean = NeuralNetwork()
11 autoenc_clean.load(H4P2C1_NN)
12
13 autoenc_noise = NeuralNetwork()
14 autoenc_noise.load(H4P2C2_NN)
15
16 # Plot training error vs epoch
17 train_errors = autoenc_clean.train_errors, autoenc_noise.train_errors
18 epochs = 10*np.arange(len(train_errors[0])), 10*np.arange(len(train_errors[1]))
19 fig1, ax1 = plt.subplots(1,2,figsize=(12,6))
20 for i in range(2):
21     ax1[i].plot(epochs[i], train_errors[i])
22     ax1[i].set_xticks(np.append(np.arange(0,epochs[i][-1],20),epochs[i][-1]))
23     ax1[i].set_xlabel("Epoch\n({})".format(int_to_roman(i+1)))
24     ax1[i].set_ylabel("Train Error")
25 ax1[0].set_title("Reconstructing Classifier~*$")
26 ax1[1].set_title("Denoising Classifier~$ ")
27 fig1.text(0.02, 0.01, '* Training early stopped at epoch {},'
28         'then restored to epoch {}, when error is {:.3f}.\n'
29         'Training early stopped at epoch {},'
30         'then restored to epoch {}, when error is {:.3f}.\n'
31         .format(epochs[0][-1], epochs[0][-1-PATIENCE], train_errors[0][-1-PATIENCE],
32               epochs[1][-1], epochs[1][-1-PATIENCE], train_errors[1][-1-PATIENCE],
33               ), ha='left')
34 fig1.tight_layout(rect=[0, 0.08, 1, 0.95])
35 fig1.savefig(H4P2_TRAIN_PLOT)
36
37 # Plot confusion matrix
38 classifier = NeuralNetwork()
39 classifier.load(H3P1_NN)
40
41 cm = [[0 for _ in range(3)] for _ in range(2)]
42 cm[0][0] = autoenc_clean.get_cm(train_db['x'], train_db['y'])
43 cm[1][0] = autoenc_clean.get_cm(test_db['x'], test_db['y'])
44 cm[0][1] = autoenc_noise.get_cm(train_db['x'], train_db['y'])
45 cm[1][1] = autoenc_noise.get_cm(test_db['x'], test_db['y'])
46 cm[0][2] = classifier.get_cm(train_db['x'], train_db['y'])
47 cm[1][2] = classifier.get_cm(test_db['x'], test_db['y'])
48
49 errors = [[0 for _ in range(3)] for _ in range(2)]
50 errors[0][0] = autoenc_clean.classify_test(train_db['x'], train_db['y'])
51 errors[1][0] = autoenc_clean.classify_test(test_db['x'], test_db['y'])
52 errors[0][1] = autoenc_noise.classify_test(train_db['x'], train_db['y'])
53 errors[1][1] = autoenc_noise.classify_test(test_db['x'], test_db['y'])
54 errors[0][2] = classifier.classify_test(train_db['x'], train_db['y'])
55 errors[1][2] = classifier.classify_test(test_db['x'], test_db['y'])
56
57 fig2, ax2 = plt.subplots(2,3, figsize=(18,12))
58 for m in range(2):
59     for n in range(3):
60         ax2[m,n].imshow(cm[m][n], cmap='Blues')
61         ax2[m,n].set_xticks(CLASSES)
62         ax2[m,n].set_yticks(CLASSES)
63         ax2[m,n].set_xticklabels(CLASSES)
64         ax2[m,n].set_yticklabels(CLASSES)
65         ax2[m,n].tick_params(axis='both', which='both', length=0)
66         for i in range(len(CLASSES)):
67             for j in range(len(CLASSES)):
68                 c = 'w' if cm[m][n][i,j]>=50 else 'k'
69                 text = ax2[m,n].text(j, i, int(cm[m][n][i, j]), ha="center", va="center", color=c, fontsize=12)
70         ax2[m,n].set_xlabel("True Class\n({})".format(int_to_roman(n+m*2+1)), fontsize=14)
71         ax2[m,n].set_ylabel("Predicted Class", fontsize=14)
72         for num in CLASSES:
73             ax2[m,n].axvline(num-0.5, c='cornflowerblue', lw=1.5, alpha=0.3)
74             ax2[m,n].axhline(num-0.5, c='cornflowerblue', lw=1.5, alpha=0.3)
75
76 ax2[0,0].set_title("Reconstructing Classifier on Train Data\n(Overall Accuracy = {:.3f})".format(1-errors[0][0]))
77 ax2[1,0].set_title("Reconstructing Classifier on Test Data\n(Overall Accuracy = {:.3f})".format(1-errors[1][0]))
78 ax2[0,1].set_title("Denoising Classifier on Train Data\n(Overall Accuracy = {:.3f})".format(1-errors[0][1]))
79 ax2[1,1].set_title("Denoising Classifier on Test Data\n(Overall Accuracy = {:.3f})".format(1-errors[1][1]))
80 ax2[0,2].set_title("BP Classifier on Train Data\n(Overall Accuracy = {:.3f})".format(1-errors[0][2]))
81 ax2[1,2].set_title("BP Classifier on Test Data\n(Overall Accuracy = {:.3f})".format(1-errors[1][2]))
82 fig2.tight_layout(rect=[0, 0, 1, 0.96])
83
84 fig2.savefig(H4P2_CM_PLOT)
85
86 plt.show()
87 plt.close('all')
```